

Estimating the Number of Legal Chess Diagrams Using Monte Carlo Simulations

Sven Hans

October 2025

Abstract

The total number of legal chess diagrams remains unknown due to the immense combinatorial complexity of chess and the difficulty of verifying legality/ reachability. This work presents a Monte Carlo sampling approach that operates on the level of material classes—distinct configurations of chess pieces. For each class, random diagrams are generated and tested for legality using the `python-chess` library. By combining combinatorial enumeration of diagrams within each class with empirically estimated legal ratios, an overall upper bound for the number of legal chess diagrams without promotions is derived. Our results suggest an upper bound of approximately 2.3×10^{40} legal diagrams, corresponding to roughly one in every five hundred random piece placements being structurally legal. These findings provide a more granular perspective on the distribution of legality across material classes and offer a foundation for future refinements toward reachable position estimation.

1 Introduction and Motivation

1.1 Scientific Overview

Estimating the total number of possible chess diagrams or positions has long been a topic in chess mathematics. Claude Shannon [1] gave an early rough estimate of the number of possible chess positions. His statement reads:

“The number of possible positions, of the general order of $64!/(32! 8!^2 2!^6)$, or roughly 10^{43} .”

This statement originates from Shannon’s foundational 1950 paper. His work focused primarily on the complexity of chess as a game rather than the enumeration of strictly legal positions. Modern studies like Chinchalkar 1996 [2] found an upper bound for the maximum legal chess positions by basic combinatoric and applying restrictions based on chess rules. In 2015 Steinerberger [3] proposed an upper bound for legal diagrams by incorporating the color-lock restriction of bishops. He additionally conjectures that an upper bound of $\approx 10^{35}$ for legal diagrams, if one could only incorporate legal pawn structures. Additionally he suggests a Monte Carlo approach to find a tighter upper bound. 2021 Tromp [4] implemented a Monte Carlo approach and estimated with 1 Mio samples the upper bound of legal positions (he calls them *urpositions*) to $\approx 10^{44}$. It is not clear to the author, if the approach is restricted to games without promotions or not.

Due to the immense combinatorial explosion, exact enumeration remains computationally infeasible. Incorporating all chess rules in a combinatoric deduction of the upper count of legal diagrams or positions seems to be too complex to solve. Monte Carlo sampling is currently the most practical approach.

1.2 Motivation

Despite decades of research, no exact enumeration of legal chess diagrams or positions exists as of 2025. The challenge primarily comes from two aspects: the complex verification of legality (ensuring positions obey chess rules such as valid pawn placement, single kings per side, castling and en passant rights, and promotions) and reachability (ensuring the position can arise from the initial setup by a series of legal moves).

In addition, many authors apply different restrictions and assumptions when defining or estimating the number of chess positions or diagrams. These include whether promotions are allowed, whether illegal pawn placements are excluded, or whether only reachable positions are considered. Such differences lead to significantly varying numerical estimates. Table 1 summarizes typical definitions and includes representative values from the literature.

Definition	Description	Upper bounds
All diagrams	All possible placements of pieces on the board, ignoring any legality constraints.	$\sim 10^{43}$ [1]
Legal diagrams	Diagrams satisfying chess rules (e.g., exactly one king per side, no pawns on rank 1 or 8) without promotions	$\approx 1,5 \times 10^{40}$ [3] $\approx 2.9 \times 10^{39}$ [4] 4×10^{37} [6]
Legal positions	Legal diagrams with additional information including side to move, castling rights and en passant possibilities.	$\approx (1.78 \pm 0.04) \times 10^{46}$ [2] , $\approx (4.8 \pm 0.04) \times 10^{44}$ [4]
Reachable positions	Legal positions that can be reached from the initial setup through valid move sequences.	
Legal with promotions	Legal positions including all possible promotions (e.g., multiple queens).	

Table 1: Different definitions of chess *diagrams* and *positions* with corresponding approximate estimates from the literature.

In this paper we try a similar approach like Tromp [4], however we will count legal diagrams without promotions, not positions. In contrary to Tromp’s work, we will sample multiple sub spaces defined by single material classes with Monte Carlo and not the full space of diagrams at once.

2 Methods

2.1 Material Classes

We define a *material class* as a set of specific chess pieces. Each material class specifies the number of kings K , queens Q , rooks R , bishops B , knights N , and pawns P present for white and black.

2.1.1 Analytical Calculation

The number of different material classes can be computed as follows: For a single side, the total number of possible piece count combinations with no promotions (limits: $K = 1, Q \leq 1, R \leq 2, B \leq 2, N \leq 2, P \leq 8$) is:

$$\text{num_combinations} = (K) \cdot (Q + 1) \cdot (R + 1) \cdot (B + 1) \cdot (N + 1) \cdot (P + 1)$$

$$1 \cdot (1 + 1) \cdot (2 + 1) \cdot (2 + 1) \cdot (2 + 1) \cdot (8 + 1) = 486$$

Since white's and black's piece counts are independent, the total number of combined material classes is:

$$486^2 = 236,196$$

2.1.2 Python Implementation

The following Python code enumerates all material classes, checking piece count constraints and total pieces:

```
import itertools

limits = {"K": 1, "Q": 1, "R": 2, "B": 2, "N": 2, "P": 8}

def all_side_materials():
    side_classes = []
    for q in range(limits["Q"] + 1):
        for r in range(limits["R"] + 1):
            for b in range(limits["B"] + 1):
                for n in range(limits["N"] + 1):
                    for p in range(limits["P"] + 1):
                        total_pieces = 1 + q + r + b + n + p #
                            Including 1 king always
                        if total_pieces <= 16: # Max pieces per
                            side
                            side_classes.append({
                                "K": 1, "Q": q, "R": r, "B": b, "
                                    N": n, "P": p
                            })
    return side_classes

white_materials = all_side_materials()
```

```

black_materials = all_side_materials()

material_classes = [
    (w, b) for w, b in itertools.product(white_materials,
        black_materials)
]

print(len(material_classes))  # 236196

```

Listing 1: Enumeration of all material classes in chess

Note: The check for total pieces per side ≤ 16 (half the board) and total pieces ≤ 32 can be added for enhanced correctness. When including promotions, these limits on individual pieces can be increased, e.g., queen count from 0 to 9, representing promotions to multiple queens.




Figure 1: Distribution of the count of different material classes for chess pieces count of 2-32

Figure 1 shows the distribution of different material classes over the number of chess pieces.

2.2 Diagram Enumeration Within a Material Class

The total number of ways to distribute the given pieces of a material class over 64 squares (ignoring legality) can be expressed combinatorially as:

$$\text{Diagrams} = \frac{64!}{(64 - n_w - n_b)! \cdot \prod_i c_i!}$$

where:

- n_w and n_b denote the number of white and black pieces, respectively,
- c_i denotes the count of identical pieces (e.g., two white rooks, eight black pawns).

Combinatorial Derivation. Let n denote the total number of pieces to be placed on the 64 squares of the chessboard. If all pieces were distinguishable (e.g., labeled), the number of possible placements would correspond to the number of injective mappings from an n -set to the 64 squares:

$$64 \cdot 63 \cdot 62 \cdots (64 - n + 1) = \frac{64!}{(64 - n)!}.$$

In practice, however, many pieces are identical (e.g., the eight white pawns or two black rooks). Let c_i denote the count of identical pieces of type i . Since permutations among identical pieces do not yield distinct arrangements, we must divide by the product of the corresponding factorials. Thus, the total number of distinct arrangements (ignoring legality) is:

$$\text{Positions} = \frac{64!}{(64 - n)! \prod_i c_i!}.$$

This formula counts purely geometric placements of the given pieces—commonly referred to as *diagrams*—and ignores all legality constraints (e.g., pawns on the first or eighth rank, both kings in check, etc.).

Example. White: 1 King, 1 Queen, 3 Pawns ($n_w = 5$)

Black: 1 King, 2 Rooks ($n_b = 3$)

Total $n = 8$. The factorials for identical pieces are $1!, 1!, 3!, 1!, 2!$. Hence:

$$\text{Positions} = \frac{64!}{(64 - 8)! 1! 1! 3! 1! 2!} = \frac{64!}{56! \cdot 3! \cdot 2!} = 14,871,915,636,480 \approx 1.49 \times 10^{13}.$$

Note: This formula counts only board *diagrams*. A full *position* in chess additionally includes meta-information such as the side to move, castling rights, and possible en-passant targets. Therefore, the total number of possible positions is higher, depending on the material configuration. Moreover, further constraints must be applied to filter only *legal* positions (e.g., valid king placement, pawn locations, and consistency of castling and check states).

2.2.1 Legal Diagram Verification Using `python-chess`

We use the `python-chess` library [5] to validate chess diagrams, ensuring that they satisfy all basic structural legality conditions. The validation routine checks a comprehensive set of status flags to detect inconsistencies or impossible states in a given diagram. Among others, the following criteria are verified automatically:

- Exactly one king per side (no missing or duplicate kings)
- No pawns on the first or last rank
- Valid castling rights
- Valid en passant squares and consistency with pawn structure
- No king in check if it is that side's turn
- No more than two simultaneous checks (`STATUS_TOO_MANY_CHECKERS`)

- No geometrically impossible double checks (`STATUS_IMPOSSIBLE_CHECK`), e.g., two attacking pieces aligned on the same ray with the king, or en passant positions that cannot have arisen from any legal move

If none of these error conditions are triggered, the position is marked as `STATUS_VALID`, meaning it satisfies all basic requirements of internal consistency. However, this does *not* guarantee that the diagram is reachable by a legal sequence of moves from the initial position—it merely ensures that it respects the fundamental structural rules of chess.

```
import chess

def is_legal(fen):
    board = chess.Board(fen)
    return board.is_valid()
```

Listing 2: Validation of chess positions using python-chess

Important: The `python-chess` library does *not* check for bishop color complex legality (e.g., disallowing two bishops of the same color on one side in non-promotion cases). Such configurations can occur legally only through pawn promotion. Therefore, when promotions are not explicitly modeled, the valid ratios measured here tend to be slightly overestimated compared to stricter definitions. This aspect should be addressed in further work.

2.3 Monte Carlo Sampling and Legal Ratio Estimation

For each material class, we estimate the fraction of legal diagrams by Monte Carlo sampling. Random diagrams are generated by placing the pieces of the class uniformly at random on the 64 squares, and each diagram is tested for legality using the `python-chess` library (see Section 2.2.1).

Let n denote the number of sampled diagrams and l the number among them that pass the legality check. The legal ratio for the material class is then estimated as:

$$\hat{r} = \frac{l}{n} \quad (1)$$

2.3.1 Adaptive Sample Size Determination

To ensure stable and reliable estimates while avoiding unnecessary computational overhead, we employ an adaptive sampling strategy with an empirical convergence criterion:

1. **Initial sampling:** Begin with $n_0 = 1,000$ samples and compute the initial legal ratio \hat{r}_0 .
2. **Iterative refinement:** Double the sample size in each iteration ($n_{i+1} = 2n_i$) and recompute the legal ratio \hat{r}_{i+1} based on n_{i+1} samples.
3. **Convergence criterion:** Continue sampling until the relative change between consecutive estimates falls below a threshold of 10%:

$$\frac{|\hat{r}_{i+1} - \hat{r}_i|}{|\hat{r}_i|} < 0.10 \quad (2)$$

or

$$\hat{r}_{i+1} - \hat{r}_i < 10^{-6} \quad (3)$$

4. **Termination:** Once this criterion is satisfied, we accept \hat{r}_{i+1} as the final estimate for the material class.

This approach ensures that the legal ratio estimate stabilizes with increasing sample size, as additional samples no longer produce substantial changes in the estimated proportion. The doubling strategy allows the method to adapt automatically to the characteristics of each material class: classes with high variance or rare legal diagrams naturally receive more samples until convergence is achieved, while classes with more stable ratios converge with fewer samples.

2.3.2 Statistical Properties

While this procedure does not provide formal confidence intervals, the convergence criterion ensures practical stability of the estimates. For material classes where convergence was achieved, the final sample sizes ranged from approximately 1,000 to 64,000, with a median of around 4,000 samples per class.

To assess the overall precision of our final estimate $N_{\text{legal}} \approx 2.28 \times 10^{40}$, we note that the standard error of a binomial proportion is $\text{SE}(\hat{r}) = \sqrt{\hat{r}(1 - \hat{r})/n}$. Given the sample sizes used, typical relative standard errors for individual material classes are on the order of 1–5%. The total uncertainty in N_{legal} arises from the combination of these individual uncertainties across all material classes, weighted by their respective contributions to the total count.

Remark: Future work could employ more formal statistical methods such as sequential probability ratio tests or variance-based stopping rules with explicit confidence intervals. However, for the purpose of establishing an order-of-magnitude estimate of legal diagrams, the present approach provides sufficient reliability.

3 Results

Based on the Monte Carlo sampling approach described above, the total number of *legal chess diagrams* (i.e., piece arrangements satisfying the structural rules of chess but not necessarily reachable by legal play) was estimated.

For each material class m , the total number of diagrams was computed combinatorially, and multiplied by the empirically estimated *valid ratio* r_m obtained from the Monte Carlo simulations. The overall estimate across all material classes is therefore given by:

$$N_{\text{legal}} = \sum_{m \in \mathcal{M}} r_m \cdot N_m,$$

where:

- \mathcal{M} denotes the set of all material classes (possible distributions of pieces by type and color),
- N_m is the total number of diagrams for material class m ,
- r_m is the fraction of those diagrams that are structurally legal according to `python-chess`.

Using this procedure, the estimated upper bound for the total number of legal diagrams is:

$$N_{\text{legal}} \approx 7.562 \times 10^{39}.$$

To compute this bound we did calculate about 4000 samples per material class and checked the legality: In total about 1×10^9 sample diagrams. When compared to the total number of all possible diagrams (i.e., all unconstrained piece placements across all material classes), the global ratio of legal to total diagrams was found to be approximately:

$$\frac{N_{\text{legal}}}{N_{\text{all}}} \approx 1.92 \times 10^{-3}.$$

This indicates that fewer than one in every five hundred random piece configurations satisfies even the basic structural rules of chess, stated above. The valid ratio r_m is, however, *not constant* across material classes: positions with a small number of pieces (e.g., king and pawns only) or heavily asymmetrical material distributions exhibit a higher probability of structural legality, while classes with many pieces or numerous pawns show significantly lower ratios due to increased likelihood of overlapping or invalid pawn placements.

These results represent an upper bound, as the legality checks performed do not include color-complex restrictions for bishops or reachability constraints based on move sequences. The legality check does also not include unreachable pawn patterns, like more than 4 pawns in one column or unreachable side-pawn patterns [7].




Figure 2: Count of diagrams and legal diagrams over the rank of each material class. The rank of a classes is defined by sorting by the number of all diagrams of the classes.

In 2 one may observe, that only very few of material classes have a huge impact on the total number of possible (legal) diagrams. Therefore it would be sufficient to only analyze these classes to come up with a good approximation of the stated upper bound for legal diagrams. The noise in the legal diagrams line, is probably due to the statistical approach and an insufficient sample size. However, we also expect a variation as the true legal ratio may fluctuate between material classes with adjacent rank.

The figure 2 shows the legal ratio over all classes, with classes sorted in the same way, as in figure 2. One may observe a noisy variation of the legal ratio in dependence of the material class. There is also a systematic drift, with low ratio for classes with high rank. Meaning the fraction of illegal diagrams is higher for classes with large counts of different diagrams.

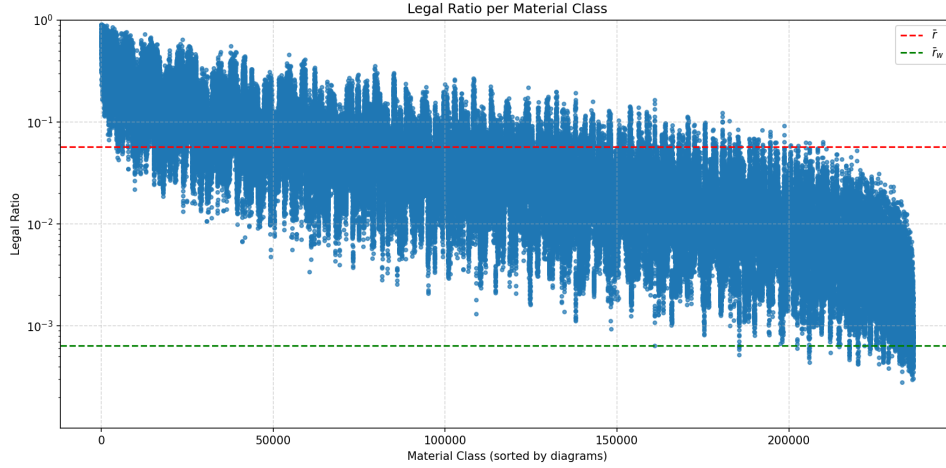


Figure 3: Legal ratio over the rank of material classes.

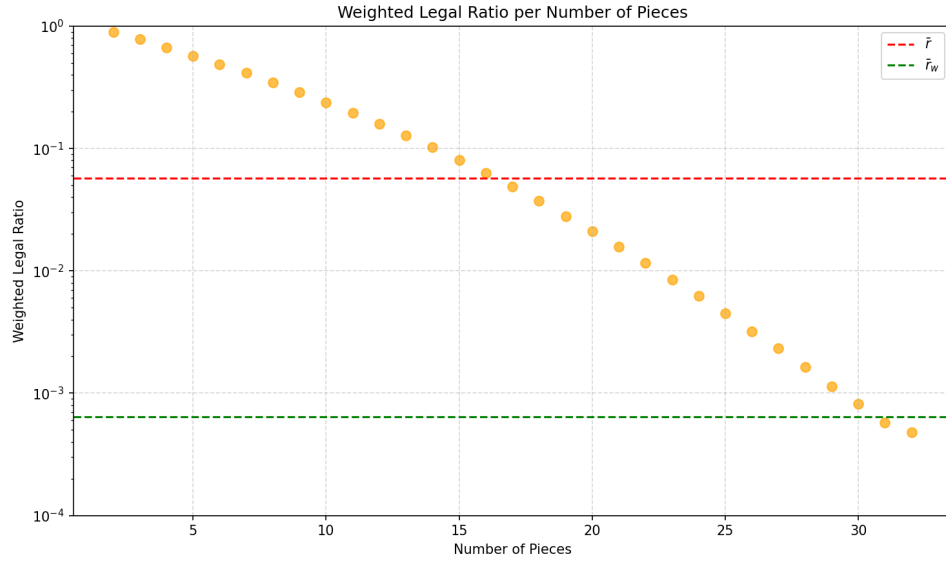


Figure 4: Average Legal ratio over number of chess pieces

Figure 4 shows the average legal ratio over the number of chess pieces. One observes, that the lowest fraction of legal diagrams can be found for 32 pieces. This indicates again, that for solving the problem of finding an upper bound for legal diagrams or positions, analyzing games with high number of pieces is most important.

Finally figure 5 shows the number of total and legal diagrams over the number of chess pieces. Interestingly we find the maximum of legal diagrams for 31 pieces and not 32, like for the total number of diagrams.

4 Summary and Conclusion

This study demonstrates that Monte Carlo simulations can provide meaningful statistical estimates for the otherwise intractable problem of enumerating legal chess diagrams. By decomposing the search space into well-defined material classes and applying legality checks on sampled diagrams, we approximate the total number of structurally valid




Figure 5: Average (legal) diagrams ratio over number of chess pieces

configurations to be on the order of 10^{40} .

The data reveal strong variations in the legality ratio across material classes and a systematic decline for positions with a high total number of pieces, especially near the full 32-piece configuration. This indicates that illegalities such as invalid pawn placements and conflicting king states dominate in dense positions.

Our findings also highlight that only a small subset of material classes significantly contributes to the global count — primarily those with near-complete sets of pieces. This observation suggests a strategy for future refinements: focus computational effort on high-impact material classes while applying statistical interpolation for the rest.

While this work excludes promotions and reachability constraints, it provides a scalable foundation for estimating the space of *legal* (but not necessarily *reachable*) chess diagrams. Extending the method to include promoted pieces, bishop color and pawn structure constraints represents a promising direction for further research.

5 References

References

- [1] C. E. Shannon, *Programming a Computer for Playing Chess*, *Philosophical Magazine*, Vol. 41, pp. 256–275, 1950.
- [2] S.,S. Chinchalkar, *An Upper Bound for the Number of Reachable Positions*, *ICCA Journal*, Vol. 19, No. 3 (Sept. 1996), pp. 181–183. DOI:10.3233/ICG-1996-19305.
- [3] S.,Steinerberger, *On the number of positions in chess without promotion*, *International Journal of Game Theory*, Vol.,44, No.,3 (Aug. 2015), pp.,761-767. DOI:10.1007/s00182-014-0453-7.
- [4] J. Tromp, *Chess Position Ranking*, GitHub repository, Dec. 2021. <https://github.com/tromp/ChessPositionRanking>

- [5] N. Fiekas, “python-chess: A Chess Library for Python,” <https://python-chess.readthedocs.io/>
- [6] D. Gourion, *An Upper Bound for the Number of Chess Diagrams Without Promotion*, arXiv preprint *arXiv:2112.09386*, 2021.
- [7] C. McDonagh, *Counting Unreachable Single-Side Pawn Diagrams with Limitless Captures*, arXiv preprint *arXiv:2202.00428v1*, 2022.